

ARTIFICIAL INTELLIGENCE

June 2025 27 11:34



1.

INTRODUCTION TO AI

- *Artificial Intelligence is an exciting scientific discipline that studies how we can make computers exhibit intelligent behavior, e.g. do those things that human.*
- *Originally, computers were invented by Charles Babbage to operate on numbers following a well-defined procedure - an algorithm. Modern computers, even though the original model proposed in the 19th century, still follow the same idea of controlled computations. Thus it is possible to program a computer to do something if we need to do in order to achieve the goal.*

▪ DIFFERENCE BETWEEN WEAK AI AND STRONG AI

○ **Weak AI (Narrow AI):**

- **Definition:** AI that is designed and trained for a specific task.

- **Capabilities:** It can *simulate* human intelligence but doesn't truly understand or possess consciousness.

- **Examples:**

Voice assistants like *Siri, Alexa*

Recommendation systems (e.g., Netflix, Amazon)

Chatbots and image recognition tools

✓ Good at performing one task extremely well

✗ Cannot generalize to other tasks

○ Strong AI (Artificial General Intelligence, AGI):

- **Definition:** AI with *human-like cognitive abilities*—able to understand, learn, and apply knowledge across a wide range of tasks.

- **Capabilities:** It would have *consciousness, self-awareness, and true understanding*—not just simulating intelligence.

- **Examples:** Purely theoretical at this point—*no current system* is truly Strong AI.

✓ Can think, reason, and adapt like a human

✗ Still hypothetical and under research

▪ TURING TEST

- When speaking about AGI we need to have some way to tell if we have created a truly intelligent system. *Alan Turing* proposed a way called a *Turing Test*, which also compares a given system to something inherently intelligent - a real human being, and because any automatic comparison can be bypassed by a computer program, being is unable to distinguish between a real person and a computer system in text-based dialogue - the system is considered intelligent.

▪ Different Approaches to AI

There are two possible approaches to this problem:

Top-down Approach (Symbolic Reasoning)	Bottom-up Approach (Neural Networks)
A top-down approach models the way a person reasons to solve a problem. It involves extracting knowledge from a human being, and representing it in a computer-readable form. We also need to develop a way to model reasoning inside a computer.	A bottom-up approach models the structure of a human brain, consist neurons. Each neuron acts like a weighted average of its inputs, and problems by providing training data.

-There are also some other possible approaches to intelligence:

- **An Emergent, Synergetic or multi-agent approach** are based on the fact that complex intelligent behaviour can be obtained by an interaction of a large number of simple cybernetics, intelligence can emerge from more simple, reactive behaviour in the process of metasystem transition.
- **An Evolutionary approach, or genetic algorithm** is an optimization process based on the principles of evolution.
- **In a top-down approach**, we try to model our reasoning. Because we can follow our thoughts when we reason, we can try to formalize this process and program it inside reasoning.
 People tend to have some rules in their head that guide their decision making processes. **For example**, when a doctor is diagnosing a patient, he or she may realize that some inflammation going on inside the body. By applying a large set of rules to a specific problem a doctor may be able to come up with the final diagnosis. This approach relies heavily on knowledge representation and reasoning. Extracting knowledge from a human expert might be the most difficult part, because a doctor he or she is coming up with a particular diagnosis. Sometimes the solution just comes up in his or her head without explicit thinking. Some tasks, such as determining be at all reduced to manipulating knowledge.
- **Bottom-Up Approach**
 Alternately, we can try to model the simplest elements inside our brain – a neuron. We can construct a so-called artificial neural network inside a computer, and then examples. This process is similar to how a newborn child learns about his or her surroundings by making observations
- **Machine Learning (ML)**
 - How it works: Learns from data and patterns, not from hard-coded rules.
 - Types:**
 - Supervised learning – Learns from labeled examples (e.g., spam vs. not spam)
 - Unsupervised learning – Finds patterns in unlabeled data
 - Reinforcement learning – Learns through trial and error (like training a dog)
 - ✓ Very powerful and flexible
 - ✗ Needs lots of data

▪ **A Brief History of AI**


• **How did AI start?**

- In the **1950s**, scientists wanted computers to think like people.
- They used **rules** and **logic** to build programs.
- One big success was **expert systems** — programs that gave advice like a doctor or engineer.
- **BUT** it was hard:
- Writing all the rules was slow.
- Keeping the computer's knowledge updated was too much work.
- So, people lost interest. This was called the **AI Winter** in the **1970s**.

• **What changed?**

- Computers became **cheaper** and **faster**.
- We got **lots of data** (thanks to the internet).
- Scientists started using **neural networks** — a way for computers to **learn from examples**.
- Neural networks became really good at:
- Recognizing images (computer vision).
- Understanding speech.

• **Chess – A Cool Example**

- **Old method:** Computers guessed lots of moves and picked the best one using logic (search and rules).
- **Better method:** Computers learned from past human games (case-based reasoning).
- **Today:** AI learns by **playing with itself** and improving over time using **neural networks + reinforcement learning** (like how humans practice).
- That's how programs like **AlphaZero** can beat world champions!
-  **Other games AI learned to play:**
- **Go** (AlphaGo beat a world champ!)
- **Poker**
- **StarCraft II**
- **Dota 2**







AI FOR BEGINNER PART-2 / SYMBOLIC AI

Knowledge Representation and Expert Systems----

What is Knowledge?

- Knowledge is what we know and understand about the world.
- It's not just information we see or hear — it's what we learn and connect in our minds.
- For example, you read a book (data), understand the meaning (information), and then remember and use it in real life — that's knowledge.

DIKW Pyramid (From Data to Wisdom)

1. Data – Just raw facts.
 Example: The word “computer” printed in a book.
 ► It's just text — doesn't mean anything until someone reads it.
2. Information – When we understand what the data means.
 Example: You read “computer” and know it's a machine.
 ► Now the word has meaning.
3. Knowledge – When we connect information to what we already know.
 Example: You learn how a computer works, what it's used for, and where to buy one.
 ► It becomes part of your personal understanding of the world.
4. Wisdom – Knowing how and when to use your knowledge.
 Example: You decide when using a computer is helpful or why someone might not need one.
 ► It's smart decision-making based on knowledge.





Classifying Computer Knowledge Representations-----

Network Representations (Semantic Networks)

- Think of a mind map or a web of ideas.
- In our brain, we connect ideas like:
"Python → is a → programming language"
- A semantic network does the same thing on a computer — it shows concepts (nodes) and relationships (arrows or edges) between them.

🧠 Example:

[Python] — is —> [Untyped Language]

[Python] — invented by —> [Guido van Rossum]

[Python] — block syntax —> [Indentation]

✳️ *Object-Attribute-Value Triplets*

- Another way to store this network in a computer is to break each connection into 3 parts:
Object – Attribute – Value

📋 Example:

Object	Attribute	Value
Python	is	Untyped-Language
Python	invented-by	Guido van Rossum
Python	block-syntax	indentation
Untyped-Language	doesn't have	type definitions

- This is easy for a computer to store, search, and connect.

Hierarchical Representations – Like a Family Tree


- We humans think in hierarchies — big categories with smaller ones inside.
- Example:
 - 🐦 Canary is a Bird
 - 🐦 Bird is an Animal
- From this, we know:
 - All birds (including canaries) have wings.
 - So if something is a bird, it inherits bird properties.

Frame Representation – Like a Form or Template

- A frame is like a profile or a form that describes an object.
- It has slots, like fields in a form.
- Each slot holds values, default values, or even rules.

 Example: Frame for Python (programming language)

Slot	Value	Default Value	Range or Notes
Name	Python		
Is-A	Untyped-Language		(category it belongs to)
Variable Case		CamelCase	(default case style)
Program Length			5–5000 lines
Block Syntax	Indent		(uses indentation)

 This is like saying:

- Python is an untyped language.
- If we don't know the case style, we assume it's CamelCase by default.
- Most Python programs are between 5 to 5000 lines.
- It uses indentation for blocks.

Procedural Representations – "Knowledge as Actions"

- In this type, knowledge is stored as a set of actions or steps to take when something happens.
- It's like if-this-happens → then-do-this.

1. Production Rules – IF-THEN Statements

- These are simple rules that help us make decisions.
- Example (Doctor's Rule):
 - 👉 IF a patient has high fever OR high C-reactive protein
 - 👉 THEN they probably have inflammation
- Once we know one part is true, we can use it to conclude something else.

2. Algorithms – Step-by-Step Procedures

- Algorithms are a set of fixed steps to solve a problem.
- Like a recipe in cooking:
 - Step 1: Boil water
 - Step 2: Add pasta
 - Step 3: Cook for 10 minutes
- 💡 But in AI, algorithms aren't used much as direct knowledge – they're more for programming, not for understanding "facts."

3. Logic – Representing Universal Knowledge

- Logic is a **formal way to represent facts and reasoning**.
- It started with **Aristotle**, who tried to describe how humans think using rules.

Types of Logic Used in AI:

▶ Predicate Logic

- Used to describe facts like:
 - 👉 "All birds can fly" or "Socrates is a man"
- It's very powerful, but too complex for computers to handle fully.
- ✅ So we use simpler parts, like Horn clauses (used in Prolog).

▶ Description Logic

- Used to describe hierarchies and relationships between objects.
- Commonly used in the Semantic Web to define and connect knowledge online.

| What is an Expert System?

An **Expert System** is a computer program that **mimics the decision-making ability of a human expert** in a specific field, like medical diagnosis, engineering, or financial analysis.

Knowledge Base (like a brain full of expert facts)

- It stores **expert knowledge**—facts, rules, and information about a specific topic.
- This knowledge is **taken from real human experts** and entered into the system.
- It **doesn't change** during each use—it's like the system's memory.
- It helps the system **move from one problem to a solution**, so it's also called **dynamic system**

⚙️ **Inference Engine** (like the thinking part of the brain)

- This is the part that **thinks and makes decisions**.
- It looks at the problem, searches through the knowledge base, and decides **what rule to apply**.
- If it needs more info, it **asks questions** to the user (like a doctor asking about symptoms).
- Its job is to **find a path from the problem to the solution** by using the rules.

Introduction to Neural Networks-

Machine Learning

Neural Networks are a part of a larger discipline called Machine Learning, whose goal is to use data to train computer models that are able to solve problems. Machine Intelligence

We will consider the two most common machine learning problems:

1. Classification

✅ What it is:

In **classification**, the goal is to **put things into categories** (or classes).


2. Regression

✅ What it is:

In **regression**, the goal is to **predict a number**.

A Model of a Neuron

From biology we know that our brain consists of neural cells, each of them having multiple "inputs" (axons), and an output (dendrite). Axons and dendrites can conduct electrical signals can exhibit different degrees of conductivity (controlled by neuromediators).



Thus, the simplest mathematical model of a neuron contains several inputs X_1, \dots, X_N and an output Y , and a series of weights w_1, \dots, w_N . An output is calculated as:

$$Y = f\left(\sum_{i=1}^N X_i w_i\right)$$

▪ Introduction to Neural Networks: Perceptron

• Perceptron Kya Hota Hai?

Perceptron ek simple sa machine learning model hai jo 2 cheezon mein farq karna seekhta hai — jaise:

- Email spam hai ya nahi
 - Image mein cat hai ya dog
- (Only 2 options — binary classification)

• 📥 Input Aur 📤 Output

- Hum kuch inputs/features dete hain, jaise $[x_1, x_2, x_3, \dots]$
- Har input ka ek weight hota hai $[w_1, w_2, w_3, \dots]$
- Ye inputs aur weights multiply hoke total score dete hain


Formula

$$y(x) = f(w \cdot x)$$

- $w \cdot x = (w_1 \times x_1 + w_2 \times x_2 + \dots + w_n \times x_n)$
- $f()$ = step function: Agar result > 0 ho to output +1, warna -1

Perceptron Ka Goal

Training ka goal hota hai:

 Aise weights dhoondhna jo zyada se zyada inputs ko sahi classify karein.

Jab Galti Hoti Hai...

Agar perceptron kisi input ko galat predict karta hai to hum weights ko update karte hain taaki agli baar sahi ho.

Error ko define karte hain:

mathematica

$$E(w) = -\sum (w \cdot x_i \times t_i)$$

- Yahaan x_i = input vector
- t_i = true label (either +1 ya -1)
- Sirf unhi inputs ko include karte hain jo galat classify huye

Gradient Descent (Update Ka Tareeka)

Weights ko update karne ka rule:

ini

$$w = w + \eta \times x_i \times t_i$$

- η = learning rate (kitna update karna hai)
- Agar galti hui hai to weights mein thoda sa change karo
- Ye process bar-bar repeat hota hai (iterations mein)

Introduction to Neural Networks & Multi-Layer Perceptron (MLP)

What' s a Neural Network?

Think of a **neural network** like a smart calculator that learns to make predictions by looking at examples. It takes some **input**, does some math, and gives an **output** (e.g., house price).

One-Layer Perceptron:

This is the simplest model. It draws a straight line (or plane) to separate two groups of points (like red dots and blue dots) based on their features. But... it only works if the data can be separated by a straight line (linearly separable).

Multi-Layer Perceptron (MLP):

To solve more complex problems, we add more layers in between input and output. These are called hidden layers. Each layer transforms the data step by step until we get a useful output.

Now, the network can:

- Handle more than 2 classes
- Predict numbers (regression)
- Separate curved/complex data shapes

Machine Learning as Math

We want to learn a function $f(x)$ that gives the correct output (y) for each input (x).

To know how good or bad our guess is, we use a loss function \mathcal{L} — a score for how wrong we are.

- For numbers \rightarrow Use squared error
- For classes \rightarrow Use accuracy or logistic loss

The better the model, the lower the loss.

Parameters: w and b

- w (weights): Think of them as knobs to control how much each input matters
- b (bias): A fixed shift to adjust the result

We call them $\theta = \langle w, b \rangle$, and we want to tune them to reduce the loss.

Gradient Descent: Learning Step by Step

Like hiking down a mountain

- Start with random values of w and b
- Check which direction reduces the loss

- Take small steps (controlled by learning rate η)
- Repeat until we reach the bottom (lowest loss)

MLP Example Math

With more layers, the function becomes:

$z1 = w1 \cdot x + b1 \rightarrow \text{Hidden layer}$ $z2 = w2 \cdot \text{activation}(z1) + b2 \rightarrow \text{Output layer}$ $f = \text{softmax}(z2) \rightarrow \text{Final prediction}$

- $\text{activation}()$ = adds non-linearity so we can handle curves (like ReLU or sigmoid)
- $\text{softmax}()$ = turns outputs into probabilities for classification

Backpropagation: Learning Backwards

To update weights correctly, we use backpropagation — a method that applies the chain rule of calculus:

- Start from the loss
- Work backwards through the layers
- Adjust each weight using the gradient

It's like following the steps of a recipe in reverse, figuring out how each ingredient (weight) affected the final taste (loss).

Neural Network Frameworks:

What Do We Need to Train Neural Networks?

To train neural networks well, we need two things:

- **Do math with tensors (big arrays of numbers)** — like addition, multiplication, applying functions like **sigmoid or softmax**.
- **Compute gradients** — this helps the model learn by improving step-by-step (called gradient descent).

Manual vs. Automatic Gradients

- With NumPy, we can do math, but we have to manually code how the model learns (called backpropagation).
- Good frameworks (like PyTorch and TensorFlow) automatically calculate gradients — much easier!

○ CPU vs. GPU

- Neural networks are heavy! So, we use **GPUs** or even **TPUs** to make training faster by **parallelizing** (doing many things at once).

🔧 Two Popular Frameworks:

Type	TensorFlow	PyTorch
Low-Level API	TensorFlow	PyTorch
High-Level API	Keras	PyTorch Lightning

- **Low-Level API** = more control, but more work (good for research)
- **High-Level API** = much easier, just build layers and train (good for projects)

You can mix both! Like:

- Build a custom layer using low-level API
- Train with high-level API

○ Which to Learn?

Choose **either PyTorch or TensorFlow** — both are good. Many people prefer PyTorch because it's easier to debug and learn.

🎯 What is Overfitting?

Overfitting = When your model is **too smart on training data**, but **bad on new data**.

Example:

Imagine you have only 5 dots to draw a line through:

- **Simple model** (left): fits the trend → good prediction.
- **Too complex model** (right): passes through every dot perfectly → bad prediction on new data.

Model	Training Error	Validation Error
Simple	5.3	5.1 

Too complex	0 🎨	20 ✖
-------------	-----	------

- **Why Overfitting Happens:**

- Too little training data
- Model is too powerful (too many parameters)
- Noisy data

- **How to Spot Overfitting:**

- Training error is low ✔
- Validation error is high ✖

- **How to Prevent Overfitting:**

1. Add more training data
2. Use a simpler model
3. Use **regularization** (like **Dropout**) — a trick to stop overfitting

Bias-Variance Tradeoff (Simple View)

Term	Meaning
Bias	Model is too simple → misses patterns (underfitting)
Variance	Model is too complex → learns noise (overfitting)

Goal: Find a balance so your model learns the real pattern, not noise.

Conclusion:

- Use frameworks like PyTorch or TensorFlow to make your life easier.
- Use high-level APIs to quickly build models.
- Understand overfitting to make your model work well on real-world data.

